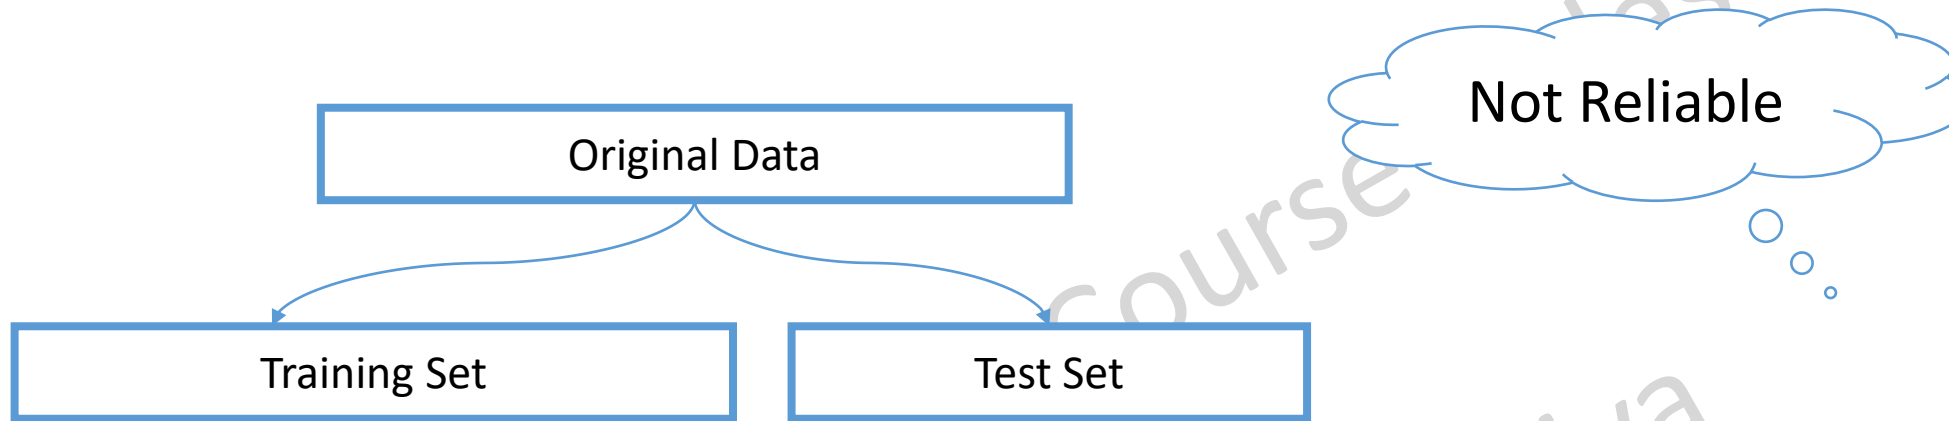


Complete Data Science and Machine Learning Using Python

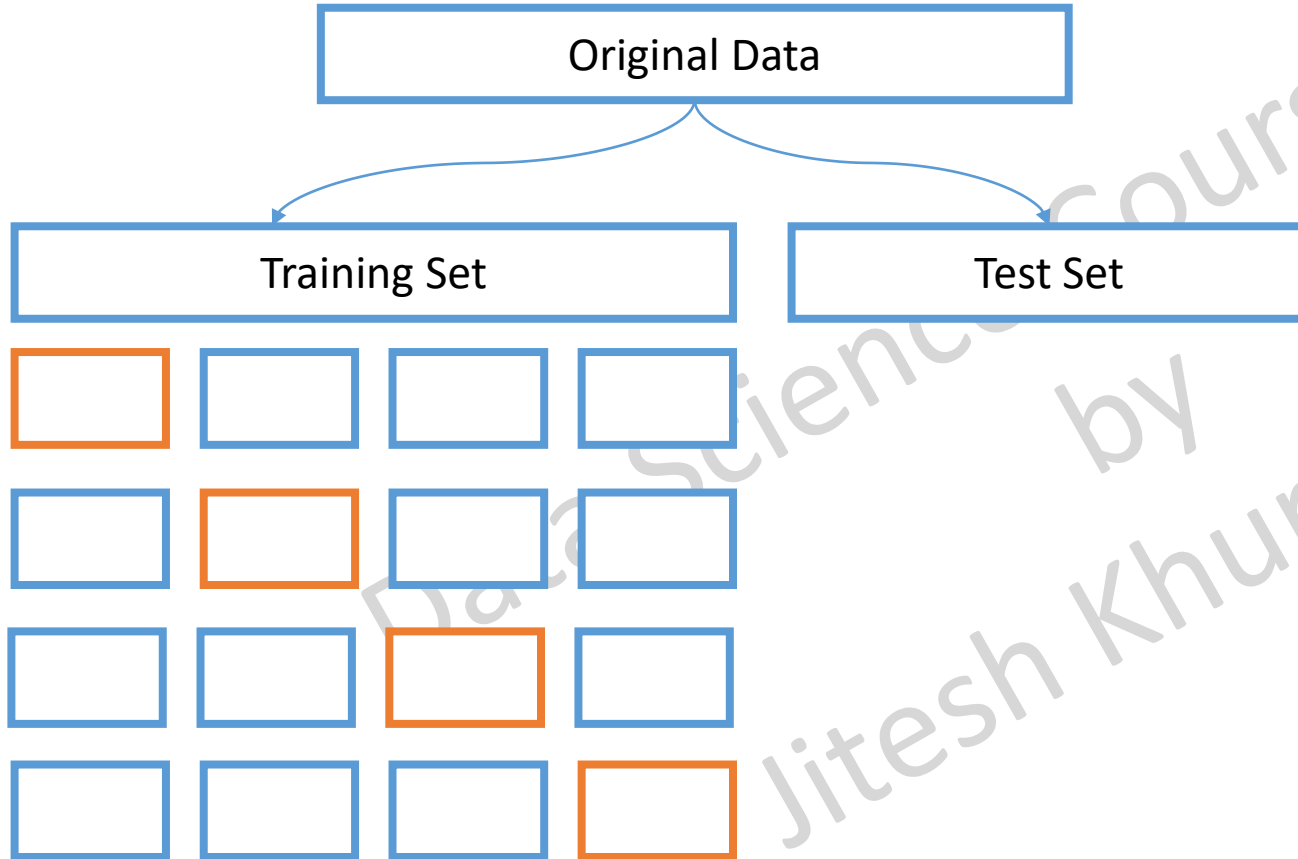
By
Jitesh Khurkhuriya

Hyperparameter Tuning

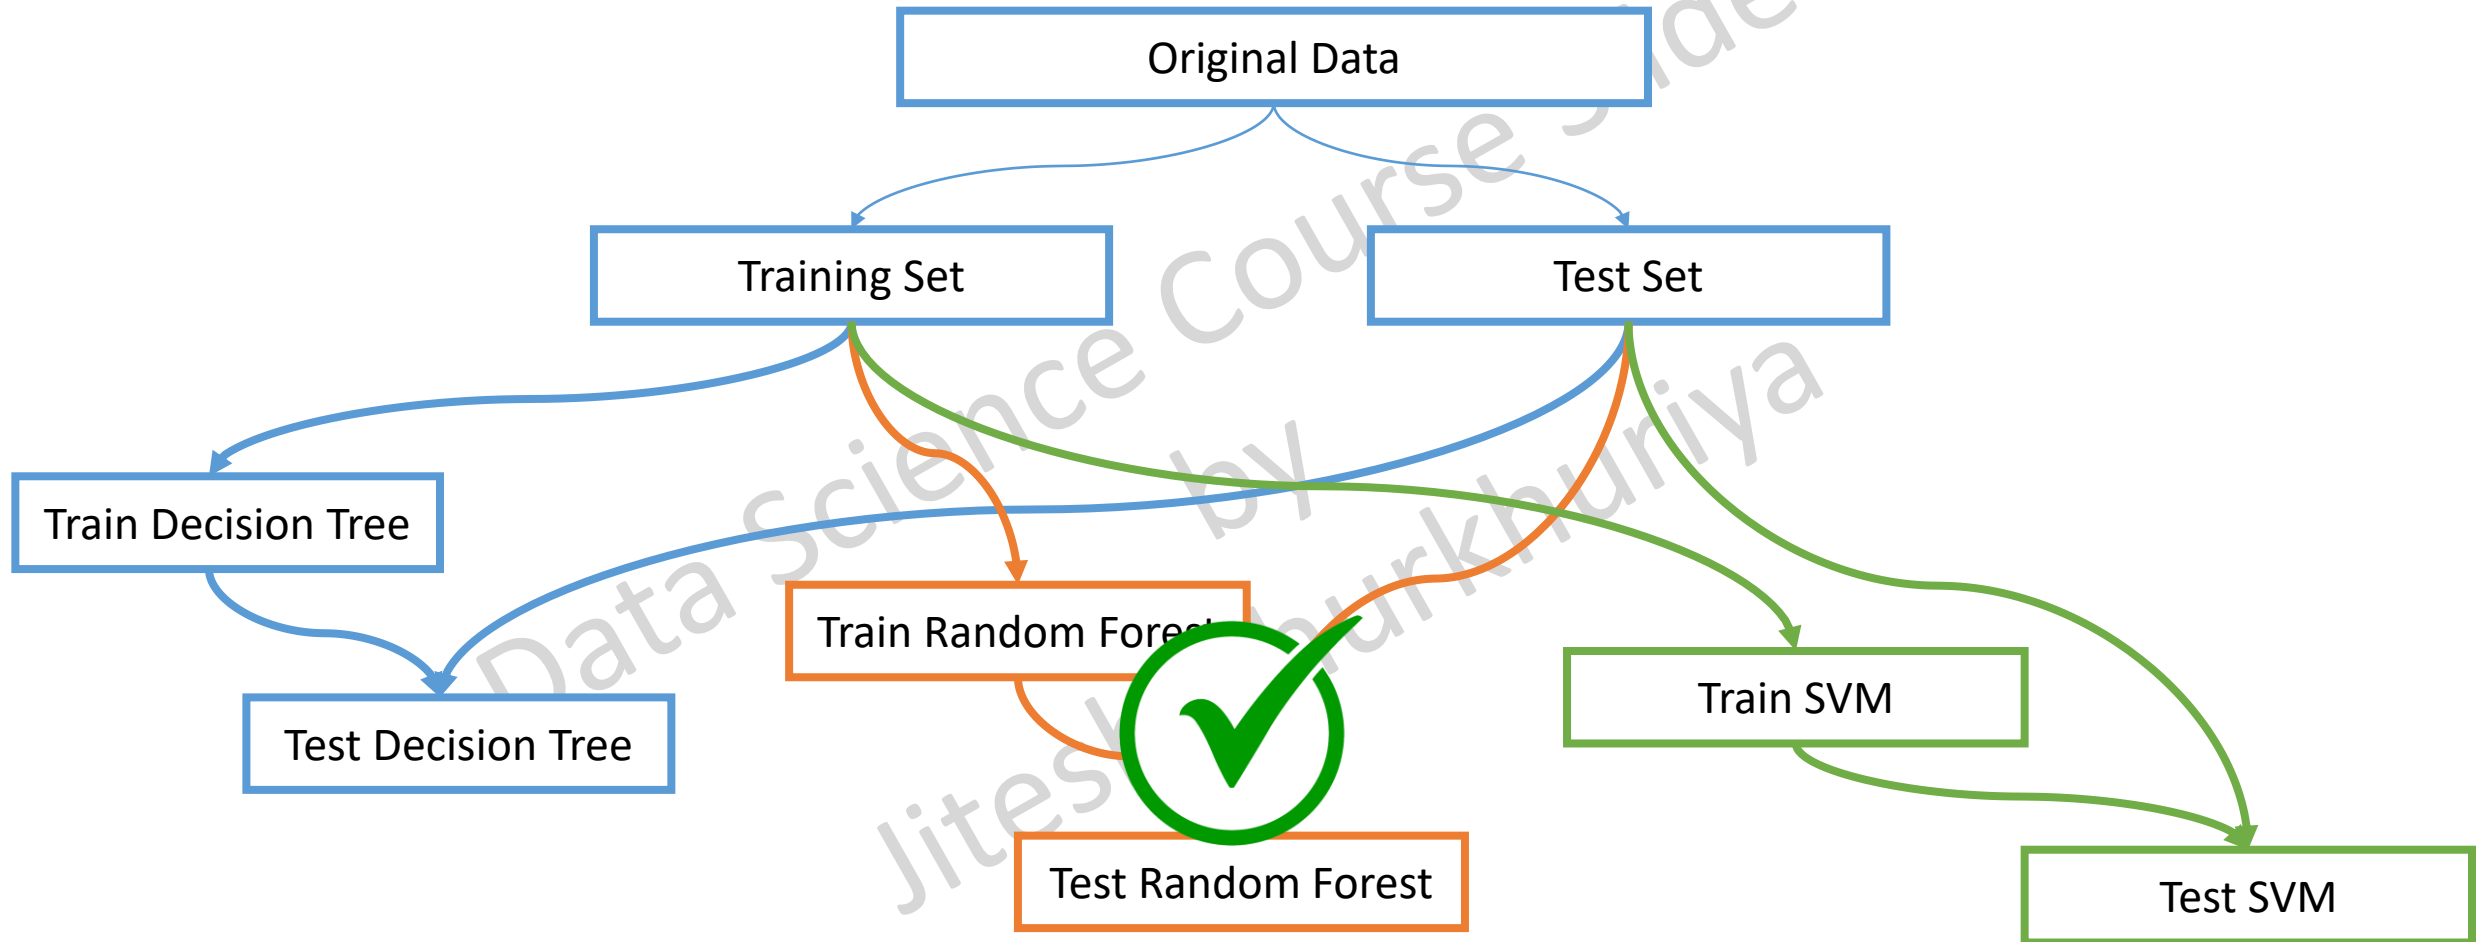
Model Selection



Cross Validation



Model Selection



Model Selection

	Decision Tree	Random Forest	Support Vector
Test :	0.7816	0.7948	0.8036
Train :	0.9044	0.8972	0.8745

Model Selection

	Decision Tree	Random Forest	Support Vector	Logistic Regression
	-----	-----	-----	-----
Test :	0.7816	0.7948	0.8036	0.8133
Train :	0.9044	0.8972	0.8745	0.8141

What are Hyperparameter and Tuning?

As per Wikipedia,

- Hyperparameter optimization or tuning is the problem of choosing a set of optimal hyperparameters for a learning algorithm.
- Hyperparameter optimization finds a tuple of hyperparameters that yields an optimal model which minimizes a predefined loss function on given independent data.
- A hyperparameter is a parameter whose value is used to control the learning process.

Cross Validation

```
19 # Import and train Decision Tree Classifier
20 from sklearn.tree import DecisionTreeClassifier
21 dtc = DecisionTreeClassifier(random_state=1234)
22
23 # Import and train Random Forest Classifier
24 from sklearn.ensemble import RandomForestClassifier
25 rfc = RandomForestClassifier(random_state=1234)
26
27 # Import and train Support Vector Classifier
28 from sklearn.svm import SVC
29 svc = SVC(kernel='rbf', gamma=0.5)
30
31 # Import and perform cross validation
32 from sklearn.model_selection import cross_validate
33 cv_results_dtc = cross_validate(dtc, X, Y, cv=10, return_train_score=True)
34 cv_results_rfc = cross_validate(rfc, X, Y, cv=10, return_train_score=True)
35 cv_results_svc = cross_validate(svc, X, Y, cv=10, return_train_score=True)
36
```

Decision Tree Classifier

- criteria
- splitter
- max_depth
- min_samples_split
- min_samples_leaf
- min_weight_fraction_leaf
- max_features
- max_leaf_nodes
- min_impurity_decrease
- min_impurity_split
- class_weight
- presort

Random Forest Classifier

- criterion
- n_estimators
- max_depth
- min_samples_split
- min_samples_leaf
- min_weight_fraction_leaf
- max_features
- max_leaf_nodes
- min_impurity_decrease
- min_impurity_split
- class_weight

Support Vector Classifier

- C – penalty parameter
- kernel
- degree
- gamma
- coef0
- shrinking
- probability
- tol
- cache_size
- class_weight
- max_iter

Logistic Regression

- penalty
- dual
- tol
- C – penalty parameter
- l1_ratio
- fit_intercept
- intercept_scaling
- class_weight
- solver
- max_iter

Random Forest Classifier

- criterion
default="gini"
- n_estimators
default=10
- max_depth
default=2
- min_samples_split
default=1
- min_samples_leaf
default="auto"
- min_weight_fraction_leaf
- max_features
- max_leaf_nodes
- min_impurity_decrease
- min_impurity_split
- class_weight

Random Forest

0.7948
0.8972

Support Vector Classifier

- C – penalty parameter
default=1.0
- kernel
default='rbf'
- degree
- gamma
gamma=0.5
- coef0
- shrinking
- probability
- tol
default=1e-3
- cache_size
- class_weight
- max_iter

Support Vector

0.8036
0.8745

What are Hyperparameter and Tuning?

As per Wikipedia,

- Hyperparameter optimization finds a tuple of hyperparameters that yields an optimal model which minimizes a predefined loss function on given independent data.



Finding the **best combination** of the **hyperparameter values** that **minimizes errors** and provides the best/optimal model.

- criterion
- n_estimators
- max_depth
- min_samples_split
- min_samples_leaf
- min_weight_fraction_leaf
- max_features
- max_leaf_nodes
- min_impurity_decrease
- min_impurity_split
- class_weight

- C – penalty parameter
- kernel
- degree
- gamma
- coef0
- shrinking
- probability
- tol
- cache_size
- class_weight
- max_iter

Hyperparameter Tuning Approaches

GridSearchCV

← Parameter 2

Parameter 1 →

	1	2	3
A	A, 1	A, 2	A, 3
B	B, 1	B, 2	B, 3
C	C, 1	C, 2	C, 3
D	D, 1	D, 2	D, 3

RandomizedSearchCV

← Parameter 2

Parameter 1 →

	1	2	3
A	A, 1	A, 2	A, 3
B	B, 1	B, 2	B, 3
C	C, 1	C, 2	C, 3
D	D, 1	D, 2	D, 3

Random Forest Classifier

criterion	n_estimators	min_samples_leaf
-----------	--------------	------------------

Support Vector Classifier

kernel	C	gamma
--------	---	-------

criteria → (gini, entropy)

n_estimators → (5, 10, 15)

min_samples_leaf → (1, 2, 5)

kernel → (rbf, linear)

C → (0.25, 0.5, 1.0)

gamma → (0.25, 0.5, 1.0)

Random Forest Classifier

criterion	n_estimators	min_samples_leaf
gini	10	1
gini	10	2
gini	10	5
gini	5	1
gini	5	2
gini	5	5
gini	15	1
gini	15	2
gini	15	5

$$2 \times 3 \times 3 = 18$$

Support Vector Classifier

kernel	C	gamma
rbf	1.0	0.5
rbf	1.0	0.25
rbf	1.0	1
rbf	0.5	0.5
rbf	0.5	0.25
rbf	0.5	1
rbf	0.25	0.5
rbf	0.25	0.25
rbf	0.25	1

$$2 \times 3 \times 3 = 18$$

Random Forest Classifier

criterion	n_estimators	min_samples_leaf
gini	10	1
gini	10	2
gini	10	5
gini	5	1
gini	5	2
gini	5	5
gini	15	1
gini	15	2
gini	15	5

$$2 \times 5 \times 5 = 50$$

Support Vector Classifier

kernel	C	gamma
rbf	1.0	0.5
rbf	1.0	0.25
rbf	1.0	1
rbf	0.5	0.5
rbf	0.5	0.25
rbf	0.5	1
rbf	0.25	0.5
rbf	0.25	0.25
rbf	0.25	1

$$4 \times 5 \times 5 = 100$$

Hyperparameter Tuning Approaches

GridSearchCV

← Parameter 2

Parameter 1 →

	1	2	3
A	A, 1	A, 2	A, 3
B	B, 1	B, 2	B, 3
C	C, 1	C, 2	C, 3
D	D, 1	D, 2	D, 3

RandomizedSearchCV

← Parameter 2

Parameter 1 →

	1	2	3
A	A, 1	A, 2	A, 3
B	B, 1	B, 2	B, 3
C	C, 1	C, 2	C, 3
D	D, 1	D, 2	D, 3

What is a Grid?

- Cartesian Product of Parameters
- Parameter 1 → 1, 2, 3
- Parameter 2 → A, B, C, D

Parameter 1 →

← Parameter 2

	1	2	3
A	A, 1	A, 2	A, 3
B	B, 1	B, 2	B, 3
C	C, 1	C, 2	C, 3
D	D, 1	D, 2	D, 3

GridSearchCV

← Parameter 2

Parameter 1 →

	1	2	3
A	A, 1	A, 2	A, 3
B	B, 1	B, 2	B, 3
C	C, 1	C, 2	C, 3
D	D, 1	D, 2	D, 3

crit	n_estimators	min_samples_leaf
gini	10	1
gini	10	2
gini	10	5
gini	5	1
gini	5	2
gini	5	5
gini	15	1
gini	15	2
gini	15	5

RandomizedSearchCV

← Parameter 2

Parameter 1 →

	1	2	3
A	A, 1	A, 2	A, 3
B	B, 1	B, 2	B, 3
C	C, 1	C, 2	C, 3
D	D, 1	D, 2	D, 3

criterion	n_estimators	min_samples_leaf
gini	10	1
gini	10	2
gini	10	5
gini	5	1
gini	5	2
gini	5	5
gini	15	1
gini	15	2
gini	15	5

Model Selection

Model Selection

Only Cross Validation

	Decision Tree	Random Forest	Support Vector	Logistic Regression
Test :	0.7816	0.7948	0.8036	0.8133
Train :	0.9044	0.8972	0.8745	0.8141

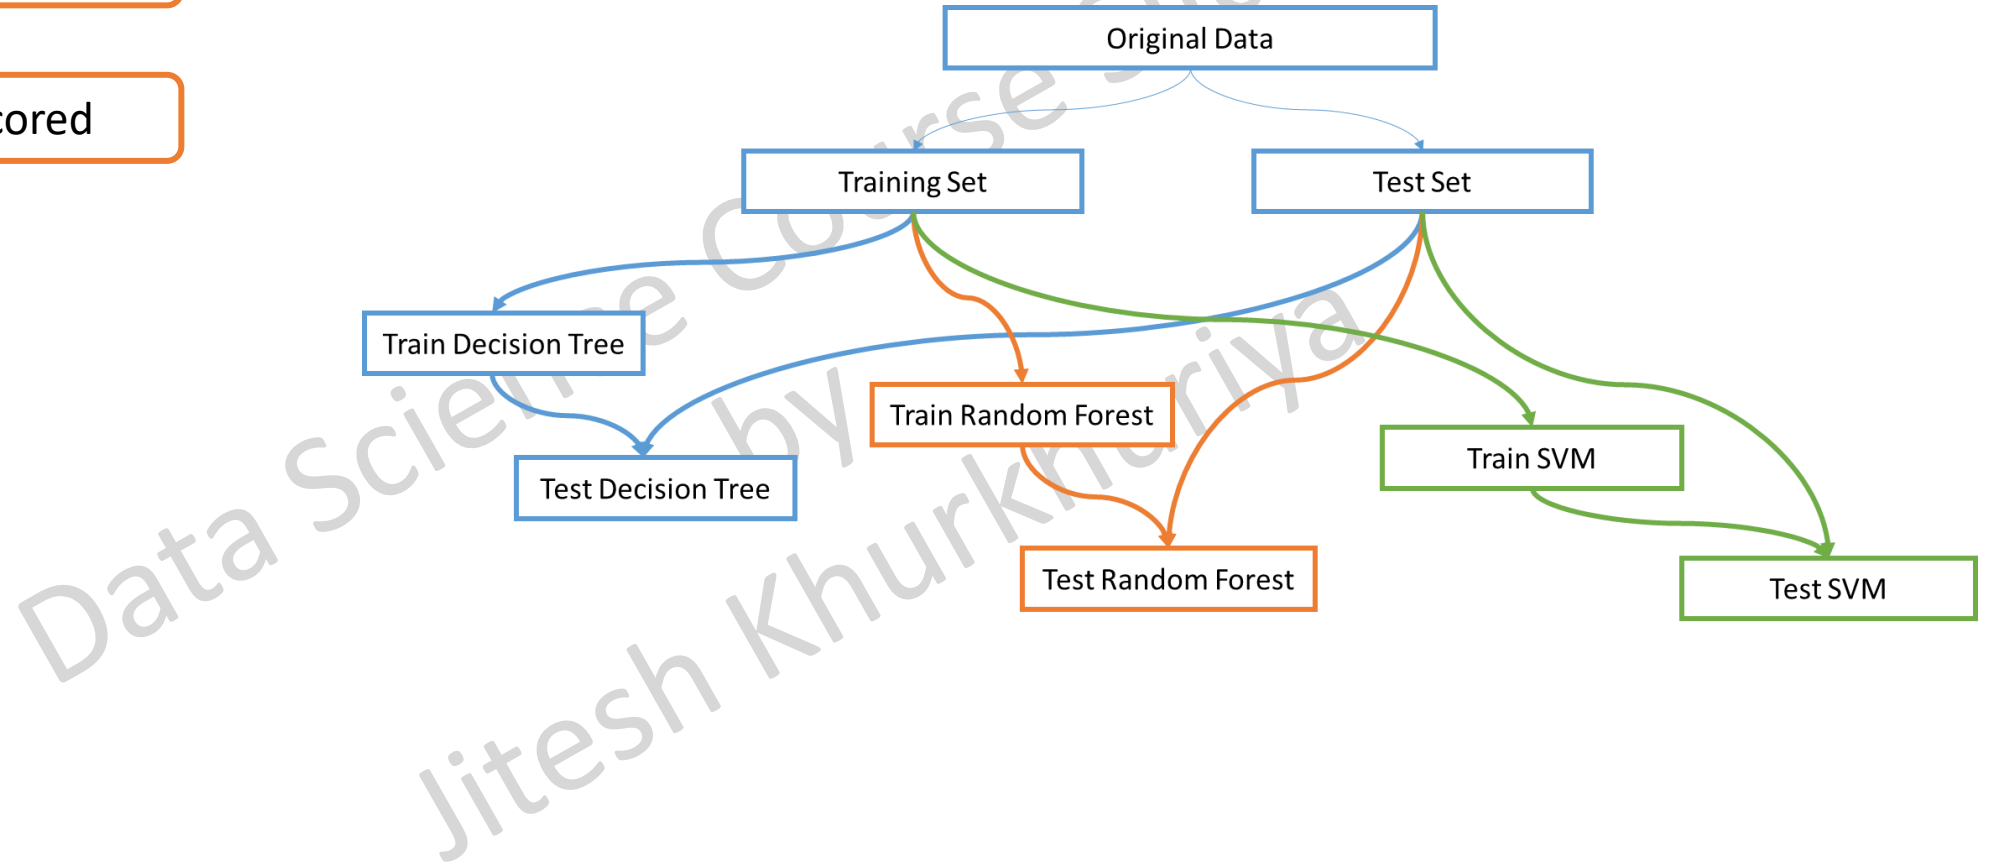
With Hyperparameter
Tuning

Test : 0.8181
Train : 0.8424

Model Selection

Train Algorithm

Select the best scored



Model Test Score for Adult Income Prediction

Train Algorithm

Select the best scored

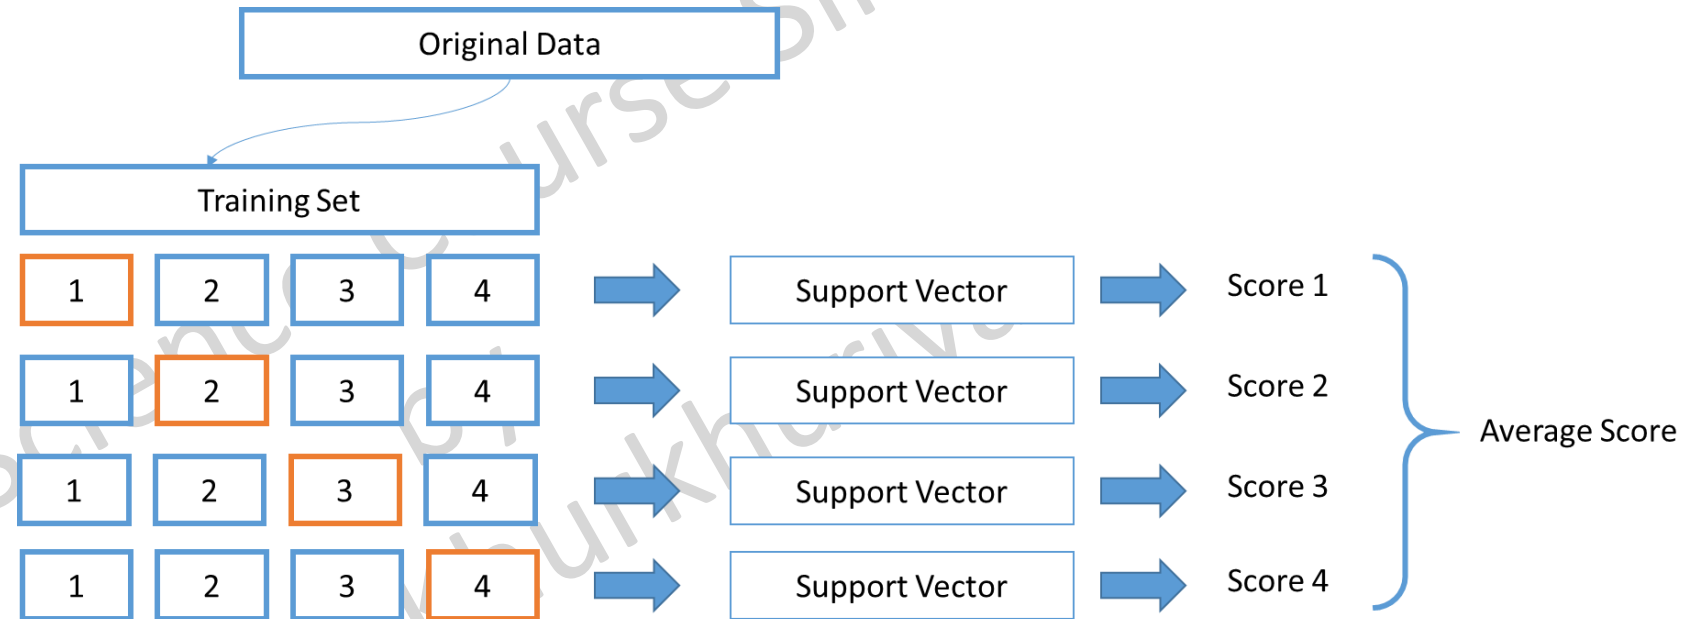
Split Random Seed	Split Size	Decision Tree	Random Forest	SVM
0	0.2	77.08%	79.18%	80.24%
123	0.2	78.39%	79.15%	80.54%
456	0.2	78.32%	78.57%	80.41%
999	0.2	76.93%	78.67%	79.73%
0	0.33	77.10%	79.30%	80.10%
123	0.33	77.81%	79.03%	79.46%
456	0.33	78.11%	79.31%	79.93%
999	0.33	77.70%	78.39%	79.49%
0	0.4	77.34%	78.96%	79.88%
123	0.4	78.44%	79.87%	79.63%
456	0.4	78.34%	79.01%	79.88%
999	0.4	77.43%	79.01%	79.79%
0	0.45	77.59%	79.30%	79.59%
123	0.45	78.06%	79.20%	79.43%
456	0.45	78.50%	79.29%	79.87%
999	0.45	77.20%	79.00%	79.71%

K-Fold Cross Validation

Train Algorithm

~~Select the best scored~~

Cross Validation



K-Fold Cross Validation

Train Algorithm

Select the best scored

Cross Validation

	Decision Tree	Random Forest	Support Vector	Logistic Regression
Test :	0.7816	0.7948	0.8036	0.8133
Train :	0.9044	0.8972	0.8745	0.8141

Train Algorithm

Select the best scored

Cross Validation

Tune Parameters

Random Forest Classifier

default="gini"

default=10

default=2

default=1

default="auto"

- criterion
- n_estimators
- max_depth
- min_samples_split
- min_samples_leaf
- min_weight_fraction_leaf
- max_features
- max_leaf_nodes
- min_impurity_decrease
- min_impurity_split
- class_weight

Random Forest

0.7948

0.8972

Support Vector Classifier

- C – penalty parameter

default=1.0

- kernel

default='rbf'

- degree

- gamma

gamma=0.5

- coef0

- shrinking

- probability

- tol

default=1e-3

- cache_size

- class_weight

- max_iter

Support Vector

0.8036

0.8745

Train Algorithm

Select the best scored

Cross Validation

Tune Parameters

```
# define parameters for Random Forest
rfc_param = {'n_estimators':[10,15,20],
             'min_samples_split':[8,16],
             'min_samples_leaf':[1,2,3,4,5]}
}
```

```
# define parameters for Support Vector Classifier
svc_param = {'C':[0.01, 0.1, 0.5, 1, 2, 5, 10],
             'kernel':['rbf', 'linear'],
             'gamma':[0.1, 0.25, 0.5, 1, 5]}
}
```

```
# define parameters for Logistic Regression
lrc_param = {'C':[0.01, 0.1, 0.5, 1, 2, 5, 10],
             'penalty':['L2'],
             'solver':['liblinear', 'lbfgs', 'saga']}
}
```

	Random Forest	Logistic Regression	Support Vector
	-----	-----	-----
Mean Test Score	: 0.8181	0.8146	0.8145
Mean Train Score	: 0.8424	0.8148	0.8429

Train Algorithm

Select the best scored

Cross Validation

Tune Parameters



cv_results_rfc - DataFrame

Index	mean_fit_time	std_fit_time	mean_score_time	std_score_time
23	0.2411	0.0330	0.0121	0.0010
18	0.1517	0.0270	0.0095	0.0034

cv_results_svc - DataFrame

Index	mean_fit_time	std_fit_time	mean_score_time	std_score_time
40	20.1890	0.3919	0.9205	0.0235
50	26.0500	0.7167	0.8961	0.0311
30	18.3952	0.5654	0.9599	0.0272

High Score Time
Higher Train-Test Score difference

	Random Forest	Logistic Regression	Support Vector
Mean Test Score	: 0.8181	0.8146	0.8145
Mean Train Score	: 0.8424	0.8148	0.8429

cv_results_lrc - DataFrame

Index	mean_fit_time	std_fit_time	mean_score_time	std_score_time
3	0.1381	0.0103	0.0021	0.0008
13	0.1808	0.0077	0.0016	0.0005

Train Algorithm

Select the best scored

Cross Validation

Tune Parameters

- Faster Predictions
- High Accuracy
- Less overfitting

cv_results_svc - DataFrame

Index	mean_fit_time	std_fit_time	mean_score_time	std_score_time
40	20.1890	0.3919	0.9205	0.0235
50	26.0500	0.7167	0.8961	0.0311
30	18.3952	0.5654	0.9599	0.0272

cv_results_rfc - DataFrame

Index	mean_fit_time	std_fit_time	mean_score_time	std_score_time
23	0.2411	0.0330	0.0121	0.0010
18	0.1517	0.0270	0.0095	0.0034

High Score Time
Higher Train-Test Score difference

	Random Forest	Logistic Regression	Support Vector
Mean Test Score	: 0.8181	0.8146	0.8145
Mean Train Score	: 0.8424	0.8148	0.8429

cv_results_lrc - DataFrame

Index	mean_fit_time	std_fit_time	mean_score_time	std_score_time
3	0.1381	0.0103	0.0021	0.0008
13	0.1808	0.0077	0.0016	0.0005



Train Algorithm

Select the best scored

Cross Validation

Tune Parameters




	Random Forest	Logistic Regression	Support Vector
Mean Test Score	0.8181	0.8146	0.8145
Mean Train Score	0.8424	0.8148	0.8429

- **Highest Accuracy**

- **High Accuracy**
- **Less overfitting**

High Score Time
Higher Overfitting

Complete Data Science and Machine Learning Using Python



Thank You!